



Introduction to Process in Computing Systems



Programs and Processes

- One way to describe the hardware of a computer system is to say that it provides a framework for executing programs and storing files.
- A *file* is a collection of data that is usually stored on disk, although some files are stored on tape.
 - UNIX treats peripherals as special files, so that terminals, printers, and other devices are accessible in the same way as disk-based files.
- A *program* is a collection of bytes representing code and data that are stored in a file.
- When a program is started, it is loaded from disk into the main memory (RAM).
- When a program is running, it is called a *process*. Most processes read and write data from files.

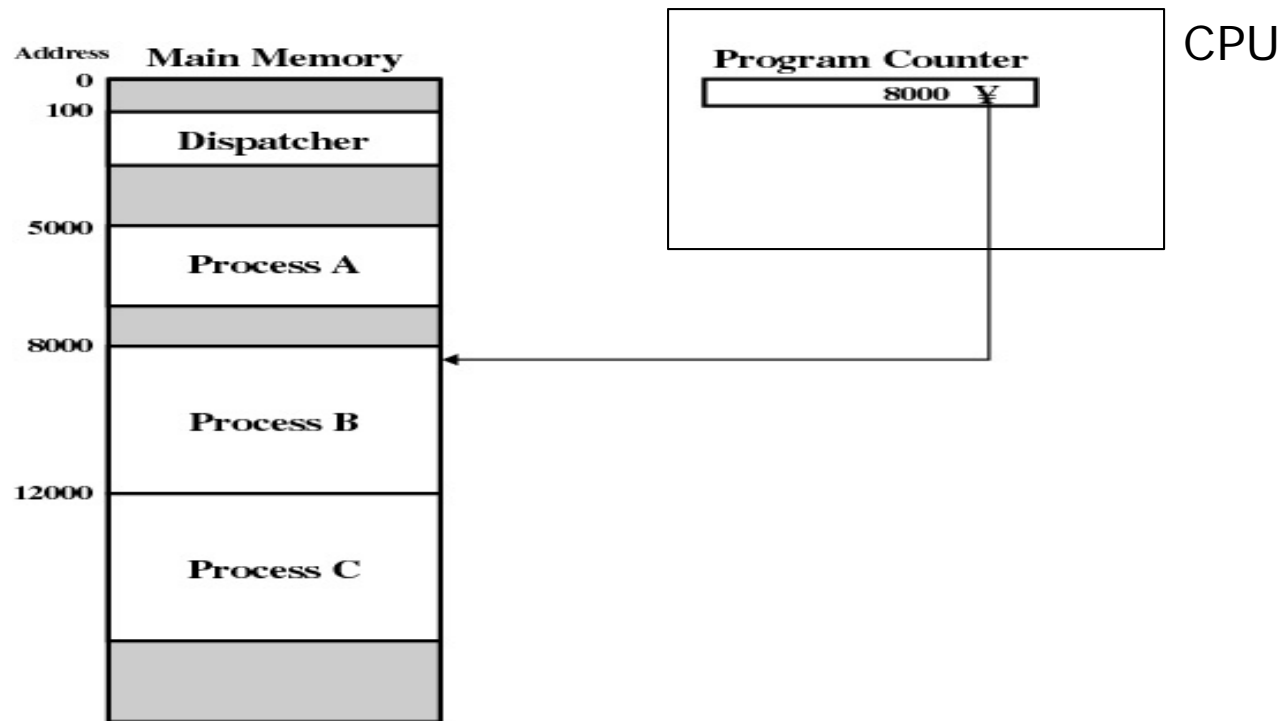


Processes

- The concept of process
 - *A program in execution*
 - The “animated spirit” of a program
 - The entity that can be assigned to and executed on a processor
- The *life of a process* is bounded by its creation and termination.
- In a typical computer system, it may be running thousands of processes at the same time.

Processes

- In order to be executed by the CPU, the code and data of a process must be located in the main memory.



**Figure 3.1 Snapshot of Example Execution (Figure 3.3)
at Instruction Cycle 13**



Show Processes in Unix

- The Unix command “ps” can display the current processes in the computer.

```
cuse93: > ps -a ... display all processes
```

PID	TT	S	TIME	COMMAND
360	console	S	0:00	/usr/lib/saf/ttymon
2434	pts/2	O	0:00	-tcsh
2431	pts/2	O	0:00	ps -a

```
cuse93: > ps -au ... display all processes with user info
```

USER	PID	%CPU	%MEM	SZ	RSS	TT	S	START	TIME	COMMAND
root	360	0.0	0.1	2987	1504	console	S	Feb 17	0:00	/usr/lib...
wlam	2434	0.2	0.1	3302	2354	pts/2	O	17:34:21	0:00	-tcsh
wlam	2431	0.1	0.2	3421	1456	pts/2	O	17:34:21	0:00	ps -au



Major Components of Operating Systems (OS)

- Process management
- Resource management
 - CPU
 - Memory
 - Device
- File system




Process Management

- One major operating system (OS) function is the sharing of limited resources among competing processes.
 - Limited resources in a typical computer system include CPUs, memory, disk space, and peripherals such as printers.
- OS shares *CPUs* among processes
 - By dividing each second of CPU time into equal-sized “slices” (typically 1/10 second) and then allocating them to processes on the basis of a priority scheme.


Process Management

- The operating system usually interleaves the execution of all processes to maximize processor utilization while providing reasonable response time.



- One component in the operating system is called *dispatcher* which takes care of the interleaving of the execution of the processes.

Trace of a Process (cont'd)



5000	8000	12000	_____
5001	8001	12001	
5002	8002	12002	
5003	8003	12003	
5004	8004 (wait for input	12004	
5005	: e.g. keyboard)	12005	
5006		12006	
5007		12007	
5008		12008	
5009		12009	
5010		12010	
5011		12011	
:		:	
(a) Trace of process A	(b) Trace of process B	(c) Trace of process C	

5000 = Starting address of program of process A
8000 = Starting address of program of process B
12000 = Starting address of program of process C

Figure 3.2 Traces of Processes of Figure 3.1



Interrupts

- A mechanism by which different processes may interrupt the normal processing of the processor.
- The classes of Interrupts
 - *Timer*: Generated by a timer within the processor.
 - *Input/Output (I/O)*: Generated by an input/output device.
- Interrupt is an interruption of the normal sequence of execution of the running process.
- After interrupt is completed, the normal program execution is resumed.
- Interrupts are provided primarily to improve processing efficiency.
 - I/O time is substantially slower than CPU processing
 - Avoid CPU waiting for slow I/O devices



1	5000	27	12004
2	5001	28	12005
3	5002		-----Time out
4	5003	29	100
5	5004	30	101
6	5005	31	102
	-----Time out	32	103
7	100	33	104
8	101	34	105
9	102	35	5006
10	103	36	5007
11	104	37	5008
12	105	38	5009
13	8000	39	5010
14	8001	40	5011
15	8002		-----Time out
16	8003	41	100
	-----I/O request	42	101
17	100	43	102
18	101	44	103
19	102	45	104
20	103	46	105
21	104	47	12006
22	105	48	12007
23	12000	49	12008
24	12001	50	12009
25	12002	51	12010
26	12003	52	12011
			-----Time out

100 = Starting address of dispatcher program

shaded areas indicate execution of dispatcher process;


first and third columns count instruction cycles;

second and fourth columns show address of instruction being executed

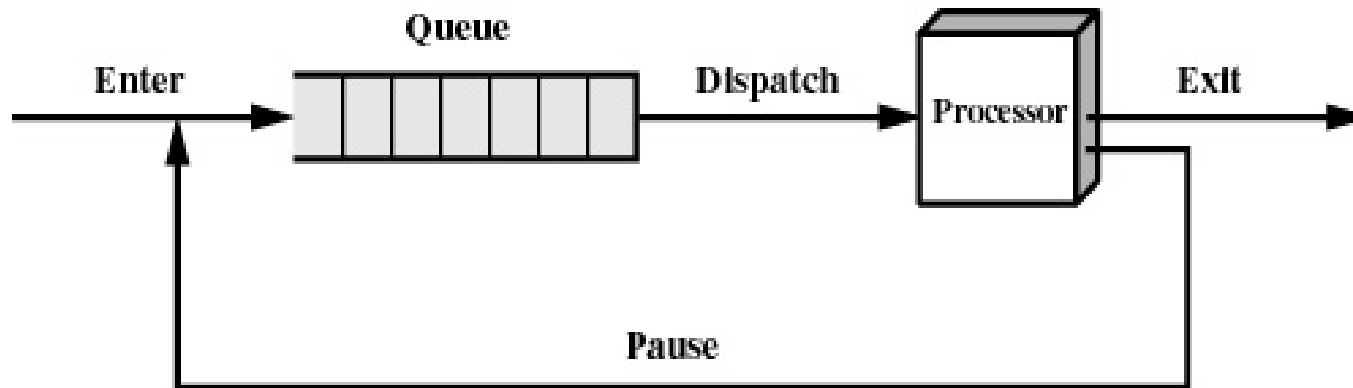
Figure 3.3 Combined Trace of Processes of Figure 3.1

A Two-State Process Model

- A *process model* is used to describe the behavior what we want the processes to exhibit.



(a) State transition diagram



(b) Queuing diagram

Two-State Process Model



A Two-State Process Model (cont'd)


- Each process must be represented in some way so that OS can keep track of it.
 - current state
 - location in memory
- Each queue item may be:
 - a pointer to a particular process
 - a data block representing a process
- *Process creation (Process Spawning):*
 - OS may create a process by itself
e.g. a user logs on to the system
 - OS may create a process on behalf of an application
e.g. requests a file to be printed
 - Application may create process
e.g. a Web browser creates a process to download a file
- *Process termination:* Termination request by application, error and fault conditions.

A Five-State Process Model





ready

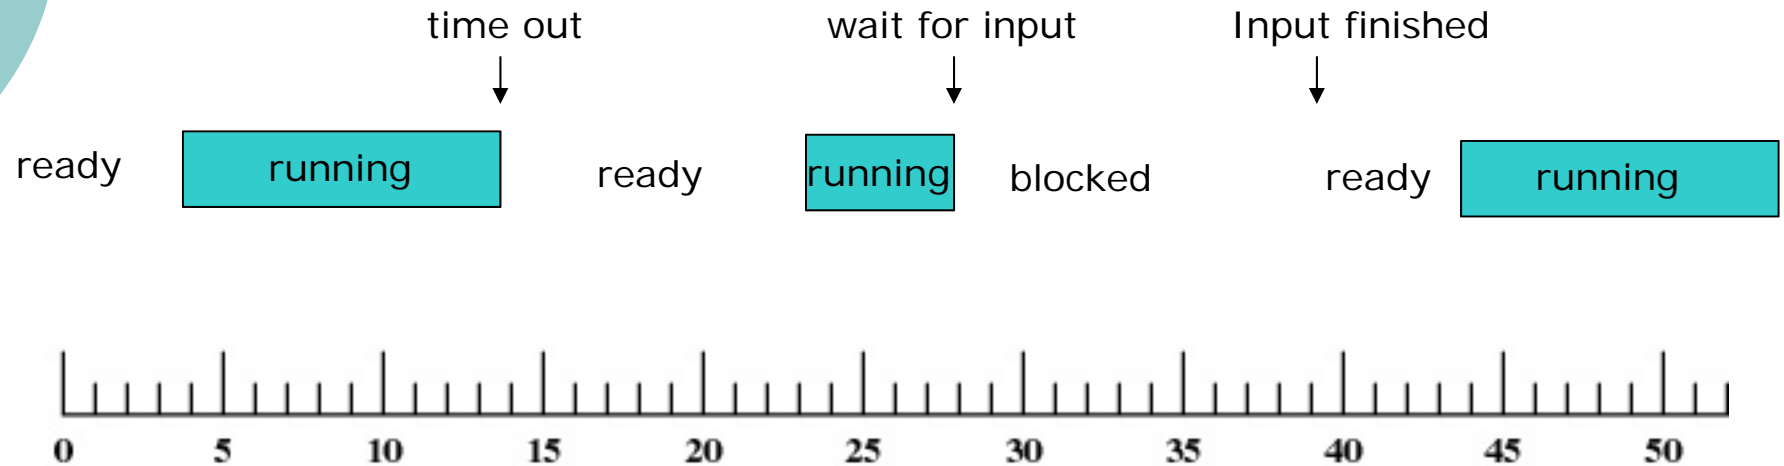
running



A Five-State Process Model



A Five-State Process Model



A Five-State Process Model

- The 2-state model does not consider that some processes in the queue may *not* be ready to execute e.g. waiting for I/O, blocked
- I/O operation is much slower than CPU computation




Figure 3.5 Five-State Process Model



A Five-State Process Model (cont'd)

- *Running*: The process is currently being executed
- *Ready*: The process is prepared to execute when given the opportunity
- *Blocked*: The process cannot execute until some event (e.g. I/O read/write) occurs
- *New*: The process has just been created but has not yet been admitted to the pool of executable processes by the OS.
- *Exit*: The process has been released from the pool of executable processes by the OS.

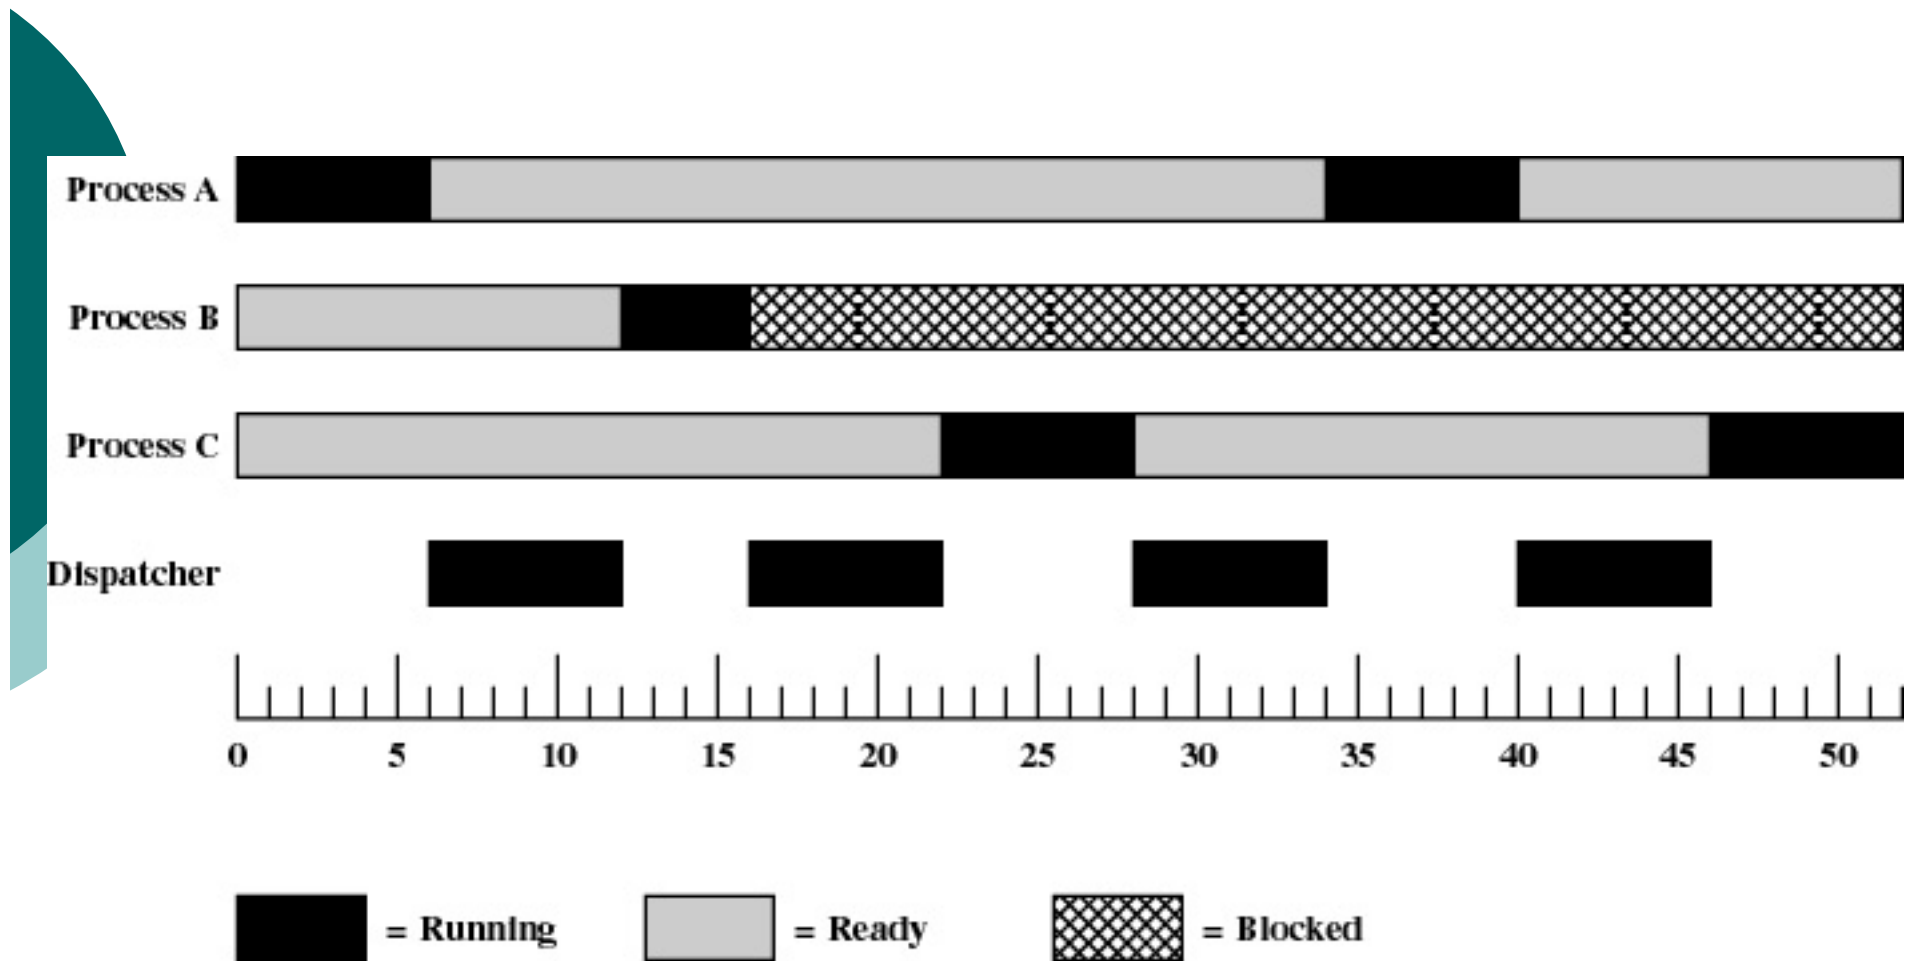
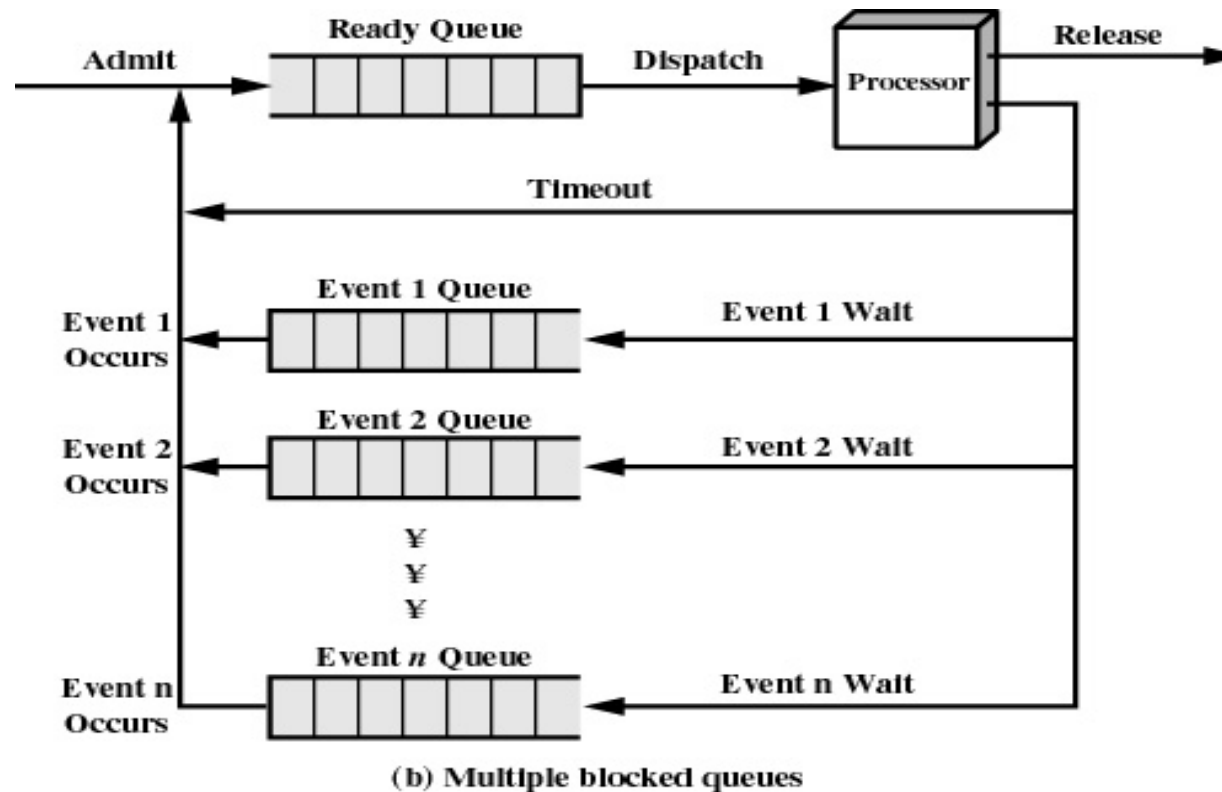



Figure 3.6 Process States for Trace of Figure 3.3

A Five-State Process Model



Queuing Model

Processes and Resources




Figure 3.9 Processes and Resources (resource allocation at one snapshot in time)

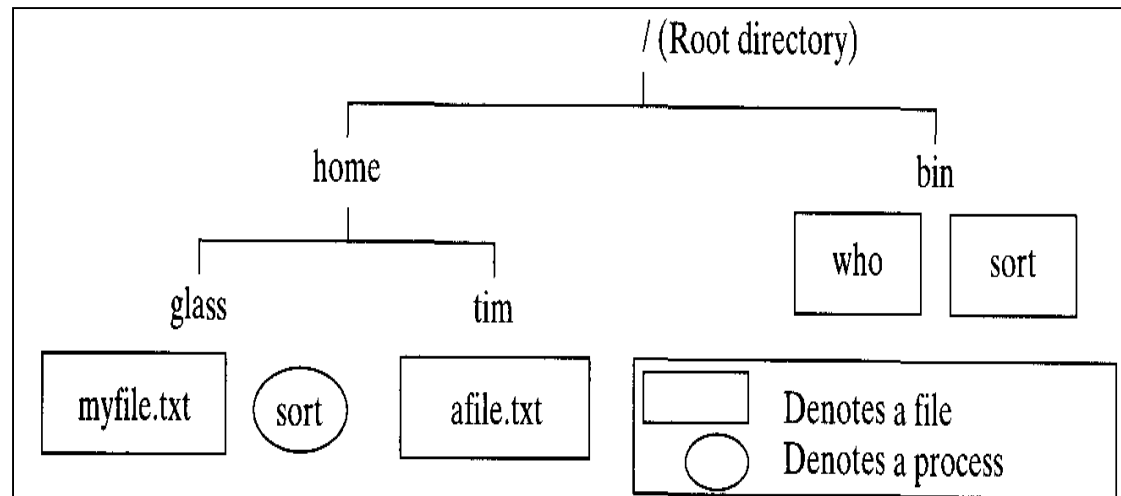


OS Functions

- OS controls events within the computer system.
- Main functions of OS:
 - schedules and dispatches processes for execution by the processor
 - allocates resources to processes
 - responds to requests by user programs for basic services.
- OS manages the use of system resources by processes

Owner of Processes and Files

- Processes and files have an *owner* and may be protected against unauthorized access.
- UNIX supports a hierarchical directory structure.
- Files and processes have a “location” within the directory hierarchy. A process may change its own location or the location of a file.
- The figure below is an illustration of a tiny UNIX directory hierarchy that contains four files and a process running the “sort” utility.





Sharing Resources

- OS shares *memory* among processes
 - By dividing the main memory up into thousands of equal-sized “chunks” of memory and then allocating them to processes
 - The chunks of a process are called *pages* and chunks of memory are called *frames*

Sharing Memory - Paging

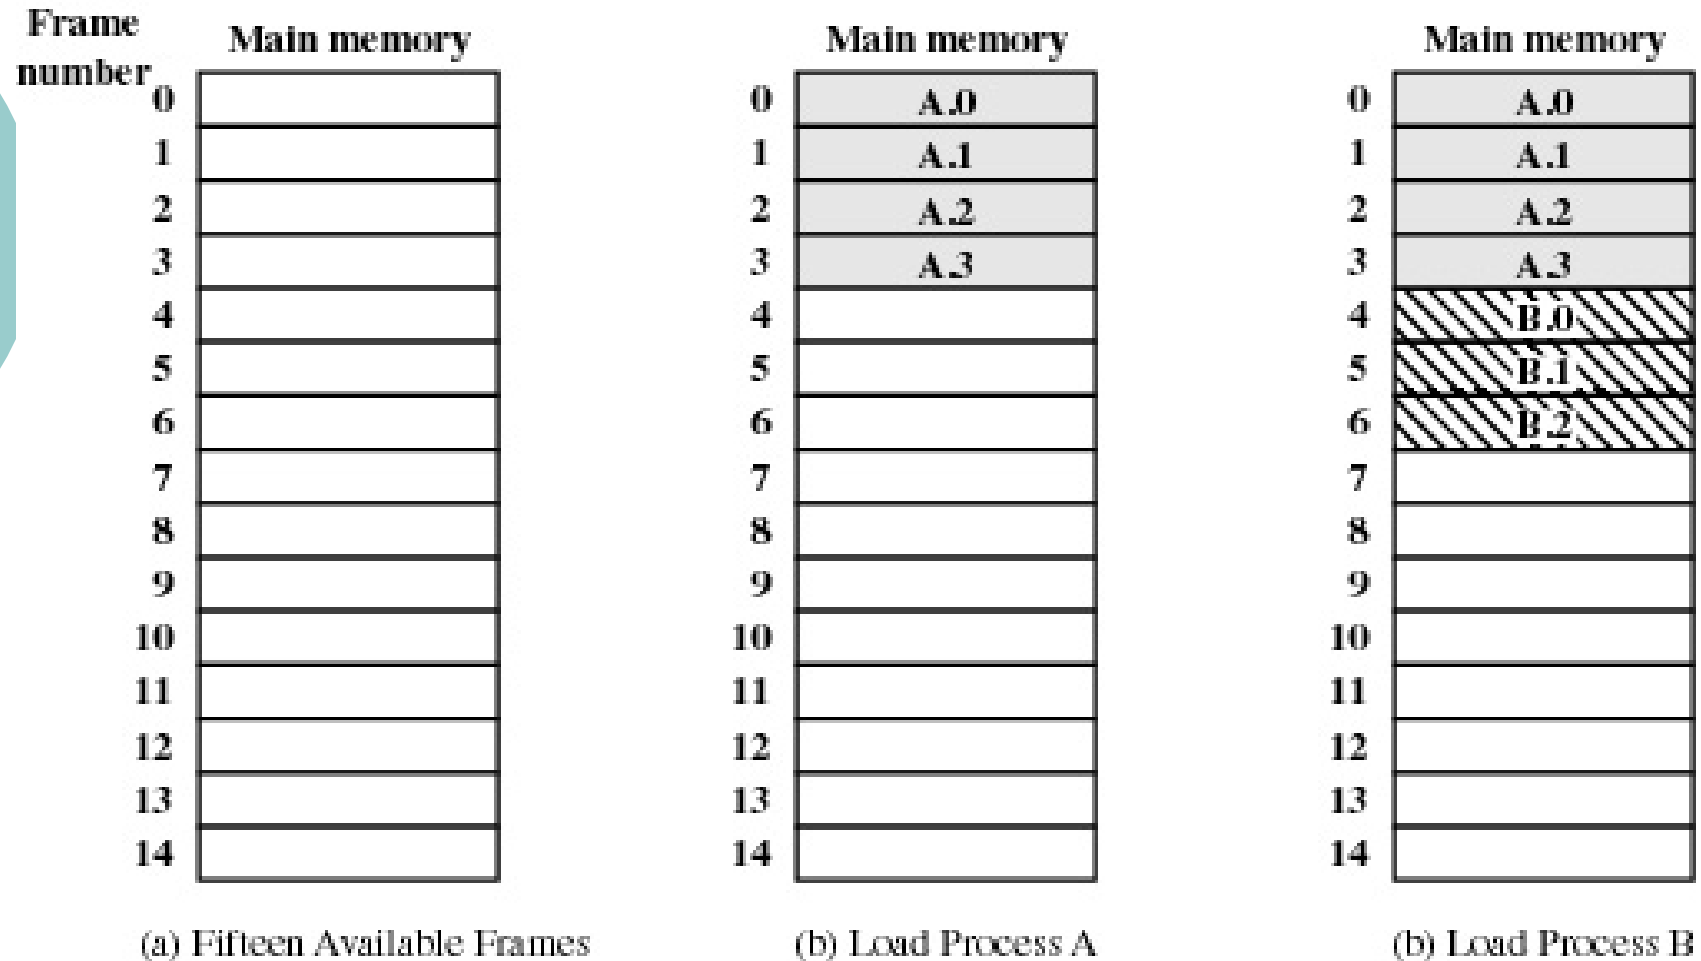


Figure 7.9 Assignment of Process Pages to Free Frames

Sharing Memory - Paging




Figure 7.9 Assignment of Process Pages to Free Frames



Sharing Memory – Page Tables for Paging

- Operating system maintains a page table for each process
 - contains the frame location for each page in the process
 - memory address consist of a page number and offset within the page

Sharing Memory – Page Tables for Paging




Figure 7.10 Data Structures for the Example of Figure 7.9 at Time Epoch (f)



Sharing Resources

- OS shares *disk space* among users
 - by dividing the disks into thousands of equal-sized “blocks” and then allocating them to users according to a quota system.
 - A single file is built out of one or more blocks.